

Assignment 5: Principal Components Analysis

Benchmark testing of alternative modeling approaches using MNIST data set from chapter 3, page 79

Based on the F1 scores of the different models and subsets of data, I would recommend using the forest classifier before the PCA transformation since the F1 score was 94%

The predictive accuracy must be weighed against the costs of the model development and implementation. Therefore, I would recommend the random forest model since it is faster. The difference in speed is not large enough to use a model with lower accuracy.

```
In [65]: import scipy.io
mnist = scipy.io.loadmat('mnist-original.mat')
X, y = mnist['data'].T, mnist['label'].T
```

Observe mnist data

```
In [66]: mnist

Out[66]: {'__globals__': [],
 '__header__': b'MATLAB 5.0 MAT-file Platform: posix, Created on: Sun Mar 30
03:19:02 2014',
 '__version__': '1.0',
 'data': array([[0, 0, 0, ..., 0, 0, 0],
                [0, 0, 0, ..., 0, 0, 0],
                [0, 0, 0, ..., 0, 0, 0],
                ...,
                [0, 0, 0, ..., 0, 0, 0],
                [0, 0, 0, ..., 0, 0, 0],
                [0, 0, 0, ..., 0, 0, 0]], dtype=uint8),
 'label': array([[0., 0., 0., ..., 9., 9., 9.]]),
 'mldata_descr_ordering': array([[array(['label'], dtype='<U5'), array(['dat
a'], dtype='<U4')]],
                                dtype=object)}
```

Review the shape of the predictor variables and the response variable

```
In [67]: X.shape
```

```
Out[67]: (70000, 784)
```

```
In [68]: y.shape
```

```
Out[68]: (70000, 1)
```

Observe a sample digit

```
In [69]: import matplotlib
import matplotlib.pyplot as plt
some_digit = X[36000]
some_digit_image = some_digit.reshape(28, 28)
plt.imshow(some_digit_image, cmap = matplotlib.cm.binary, interpolation="nearest")
plt.axis("off")
plt.show()
```



```
In [70]: # see if y can predict the correct value
y[36000]
```

```
Out[70]: array([5.])
```

Utilize the first 60,000 as a model development set and the final 10,000 as a holdout test set. Next, we develop a multiclass classifier

```
In [71]: X_train, X_test, y_train, y_test = X[: 60000], X[ 60000:], y[: 60000], y[ 60000:]
```

Begin by fitting a random forest classifier using the full set of 784 explanatory variables and the model development set of 60,000 observation. Below I walk through a stochastic Gradient Descent classifier model as an example.

```
In [72]: import time
start = time.time()
import numpy as np
shuffle_index = np.random.permutation(60000)
X_train, y_train = X_train[shuffle_index], y_train[shuffle_index]

# will use the multiclass classification method
from sklearn.linear_model import SGDClassifier
sgd_clf = SGDClassifier(random_state=42)
sgd_clf.fit(X_train, y_train)
end = time.time()
print(end-start)
```

C:\Users\lcamero\AppData\Local\Continuum\Anaconda3\lib\site-packages\sklearn\utils\validation.py:526: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

9.818561792373657

This model took 9 seconds to run. Next, I predict the response variable and score the model using the F1 score.

```
In [74]: # reshape the y train array
c,r = y_test.shape
y_test = y_test.reshape(c,)
```

```
In [115]: start = time.time()
y_train_predict = cross_val_predict(sgd_clf, X_train, y_train, cv=3)
end = time.time()
print(end-start)
```

24.580405950546265

```
In [116]: from sklearn.metrics import f1_score
```

```
In [117]: f1_score(y_train, y_train_predict, average="macro")
```

Out[117]: 0.8669597087235005

Next I fit a random forest classifier

```
In [118]: start = time.time()
from sklearn.ensemble import RandomForestClassifier
forest_clf = RandomForestClassifier(random_state=42)
forest_clf.fit(X_train, y_train)
end = time.time()
print(end-start)
```

10.599606275558472

This model took 18 seconds to complete. Next, we set up the y train predictor and use the F1 score to find the harmonic mean of precision and recall

```
In [119]: start = time.time()
y_train_predict = cross_val_predict(forest_clf, X_train, y_train, cv=3)
end = time.time()
print(end-start)
```

26.368508338928223

```
In [120]: f1_score(y_train, y_train_predict, average="macro")
```

Out[120]: 0.9397497464092263

Next use a Multiclass classifier using the K Neighbor

```
In [121]: from sklearn.neighbors import KNeighborsClassifier
knn_clf = KNeighborsClassifier()
```

```
In [122]: # create the response prediction array
from sklearn.model_selection import cross_val_predict
```

```
In [123]: start = time.time()
y_test_predict = cross_val_predict(knn_clf, X_test, y_test, cv=3)
end = time.time()
print(end-start)
```

217.35143160820007

```
In [124]: f1_score(y_test, y_test_predict, average="macro")
```

Out[124]: 0.9357079587489915

Compare test set performance across the two modeling approaches: 784 variable model versus the 95% PCA model.

Execute PCA on the full set of 70,000, generating PCA that represent 95 percent of the variability. It took 5 seconds.

```
In [86]: from sklearn.decomposition import PCA
pca = PCA(X)
```

```
In [87]: start = time.time()
pca = PCA(n_components=0.95)
X_reduced = pca.fit_transform(X)
end = time.time()
print(end-start)
```

17.362993240356445

```
In [88]: start = time.time()
forest_clf.fit(X_reduced, y)
end = time.time()
print(end-start)
```

C:\Users\lcamero\AppData\Local\Continuum\Anaconda3\lib\site-packages\ipykernel_main__.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
from ipykernel import kernelapp as app
```

40.56231999397278

```
In [89]: # reshape the y train array
c,r = y.shape
y = y.reshape(c,)
```

```
In [90]: start = time.time()
y_predict = cross_val_predict(forest_clf, X_reduced, y, cv=3)
end = time.time()
print(end-start)
```

66.93882846832275

```
In [91]: f1_score(y, y_predict, average="macro")
```

Out[91]: 0.881805729042257

Execute PCA on the set of 60,000 to build another random forest classifier.

```
In [92]: pca = PCA()
pca.fit(X_train)
```

Out[92]: PCA(copy=True, iterated_power='auto', n_components=None, random_state=None, svd_solver='auto', tol=0.0, whiten=False)

```
In [93]: start = time.time()
pca = PCA(n_components=0.95)
X_reduced_train = pca.fit_transform(X_train)
end = time.time()
print(end-start)
```

12.896737575531006

```
In [94]: start = time.time()
forest_clf.fit(X_reduced_train, y_train)
end = time.time()
print(end-start)
```

C:\Users\lcamero\AppData\Local\Continuum\Anaconda3\lib\site-packages\ipykernel_main__.py:2: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
from ipykernel import kernelapp as app
```

36.78110384941101

```
In [105]: start = time.time()
y_train_predict = cross_val_predict(forest_clf, X_reduced_train, y_train, cv=3)
end = time.time()
print(end-start)
```

62.59558033943176

```
In [106]: y_train.shape
```

Out[106]: (60000,)

```
In [107]: y_train_predict.shape
```

Out[107]: (60000,)

```
In [112]: # reshape the y train array
# c,r = y_train.shape
# y_train = y_train.reshape(c,)
```

```
In [114]: f1_score(y_train, y_train_predict, average="macro")
```

Out[114]: 0.8791328511686146

The flaw is that we did not use the test dataset. We refit the models using the test datasets and rescore. First we set up the random forest on the test dataset.

```
In [125]: start = time.time()
from sklearn.ensemble import RandomForestClassifier
forest_clf = RandomForestClassifier(random_state=42)
forest_clf.fit(X_test, y_test)
end = time.time()
print(end-start)
```

1.9611120223999023

```
In [126]: start = time.time()
y_test_predict = cross_val_predict(forest_clf, X_test, y_test, cv=3)
end = time.time()
print(end-start)
```

4.1352362632751465

```
In [127]: f1_score(y_test, y_test_predict, average="macro")
```

Out[127]: 0.8906431987279889

Set up the PCA model on the test dataset.

```
In [130]: pca = PCA()
pca.fit(X_test)
```

Out[130]: PCA(copy=True, iterated_power='auto', n_components=None, random_state=None, svd_solver='auto', tol=0.0, whiten=False)

```
In [128]: start = time.time()
pca = PCA(n_components=0.95)
X_reduced_test = pca.fit_transform(X_test)
end = time.time()
print(end-start)
```

3.0351736545562744

```
In [129]: start = time.time()
forest_clf.fit(X_reduced_test, y_test)
end = time.time()
print(end-start)
```

3.583204746246338

```
In [131]: start = time.time()
y_test_predict = cross_val_predict(forest_clf, X_reduced_test, y_test, cv=3)
end = time.time()
print(end-start)
```

7.81244683265686

```
In [132]: f1_score(y_test, y_test_predict, average="macro")
```

Out[132]: 0.7674240306924224